

CASE STUDY 3

Classification/Clustering of Emails

Balaji Avvaru, Apurv Mittal, Ravi Sivaraman

Quantifying The World | SMU | DR. SLATER

Abstract

The objective of this study is to classify emails as spam or not spam. In addition to that, cluster emails to observe any hidden group characteristics. The emails have several metadata, email body text, and sometimes attachments. The texts in the emails are parsed and analyzed for clustering and for the classification of spam emails.

Introduction

The dataset consists of five directories, two of which are spam, and the rest are not spam. Each directory contains several files of spam or not spam email and each email is stored as a single file. The format of the file is email header and body text.

To create the dataset, this study parses each file, gets the body text and various headers, and creates a pandas dataframe. The emails in the files are of the following types:

Email Type	Count
text/plain	7413
text/html	1193
multipart/alternative	326
multipart/signed	180
multipart/mixed	179
multipart/related	56
multipart/report	5
text/plain charset=us-ascii	1

Table 1-Email types

Each file type must be parsed differently.

text/plain

They are simple text and can be parsed as strings

text/html

The HTML text is parsed as HTML.

multipart

Multipart emails are emails that are large and sent in multiple parts. To get the entire message, each multipart of the message must be parsed and then appended to one message. Sometimes, the messages contain mixed multipart messages, a combination of *HTML* and *plaintext* messages. In such cases, the type of each multipart needs to be checked and then parsed appropriately.

Attachments

The email headers are parsed and if *Content-Disposition* has any value attachment, then the message has attachment, else none.

Target

All the email messages are parsed and then stored in the dataframe. The target is the message is spam or not (1 spam/0 not-spam) is computed based on folder name contains text “spam”, is treated as spam, else not spam.

Column	Non-Null	Count	Dtype
data	9353	non-null	object
target	9353	non-null	int64
Attachments	9353	non-null	int64

Table 2 - Dataframe fields

2. Methods

The email texts are parsed and stored as lower text to catch spammers who use either mixed texts or upper case to defeat spam filters. Converting all of them to lower case does lose information but reduces the number of features and is better at catching based on texts. However, sometimes mixed cases or unnatural all lower case or upper case may signify spam.

Missing Data & Imputation

There are no missing data, and there was no imputation done.

Target

Spam or not spam is the target variable, and the distribution of spam vs not-spam is as follows; the graph indicates the target variable is heavily imbalanced.




Figure 1- Count of Spam (1) vs Non-Spam (0)

The graph below shows the distribution of attachments in Spam/Not-Spam emails. The attachments are imbalanced, the non-spam emails contain proportionately higher attachments compared to spam emails. The KMeans clustering and Naïve Bayes classification doesn't required the data to be balanced.

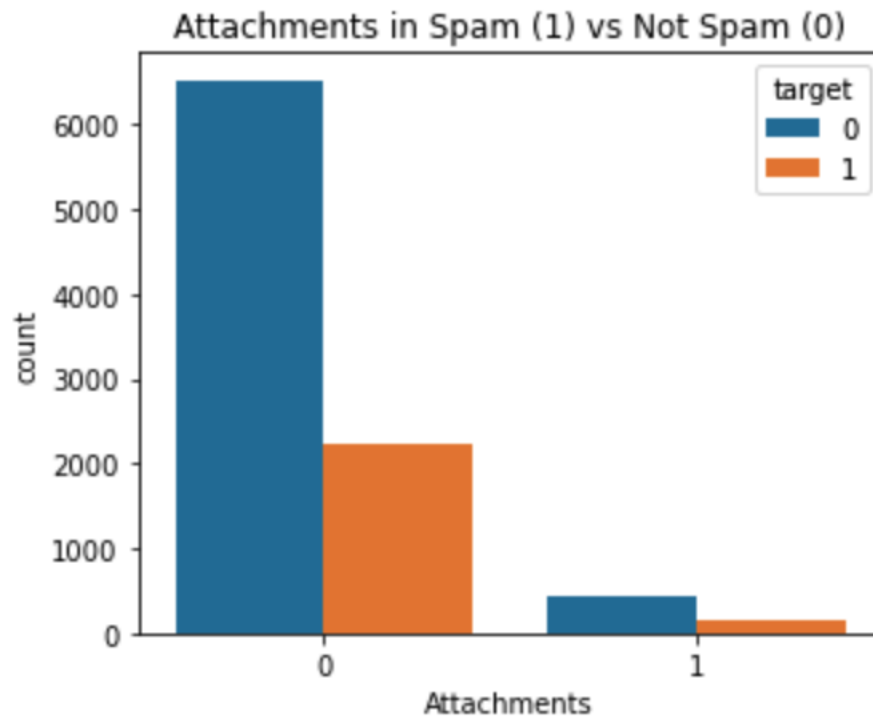


Figure 2 - Attachments in Spam vs Non-Spam (target- 0: non-spam/1: spam)

Vectorization

The body of email texts may contain a lot of smaller sized, but very common words like the, of, etc. are stripped. Some common email texts like spamassassin (part of email footer, for example) are also removed.

The extra stop words that are added to corpus including:

- Spamassassin
- Email
- Message
- nbsp
- font
- exhm
- subject
- list
- URL
- net
- http

This model uses a `Tfidf` vectorizer, to eliminate the most common words.

Clustering Models

KMeans Clustering

K-means clustering is a type of unsupervised learning, which is used when you have unlabeled data (i.e., data without defined categories or groups). The goal of this algorithm is to find groups in the data, with the number of groups represented by the variable k . The algorithm works iteratively to assign each data point to one of the k groups based on the features that are provided. Data points are clustered based on feature similarity.

The graph below shows the sum of squares at various values of k (clusters), with elbow at 15, which indicates there may be 15 optimal clusters, which this study will use.

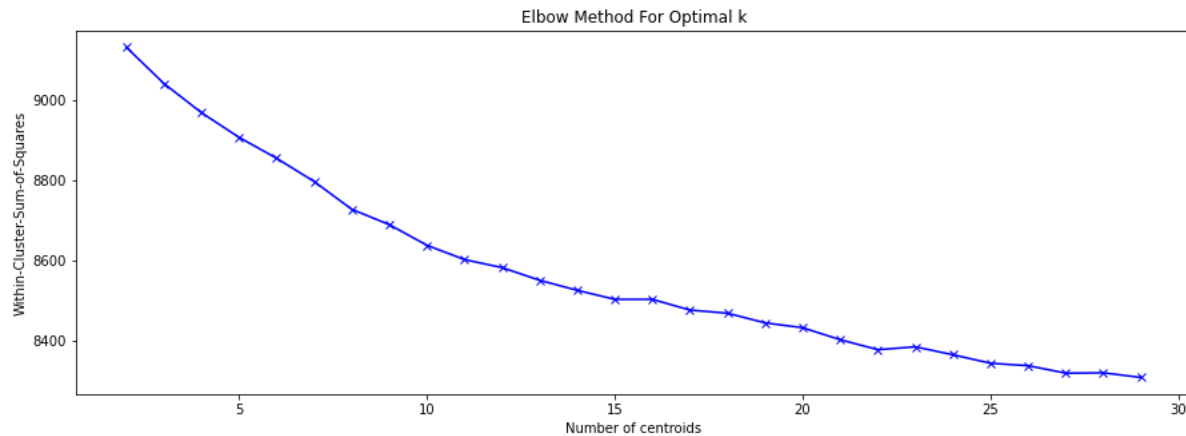


Figure 2-KMeans Range of k Values (Elbow at 15)

Visualize Clusters

This study uses a technique called t-SNE (t-distributed Stochastic Neighbor Embedding) to generate a 2-dimensional representation of this dataset, which provides an intuitive view of clusters.




Figure 3-Two-dimensional cluster of emails

Word Cloud

Each picture is a cluster, two sample clusters are shown here below. The cluster shows the emails are grouped that are related by texts (or theme of texts). The texts in the same cluster means the KMeans distance between them are low compared to the texts that are not in same cluster.

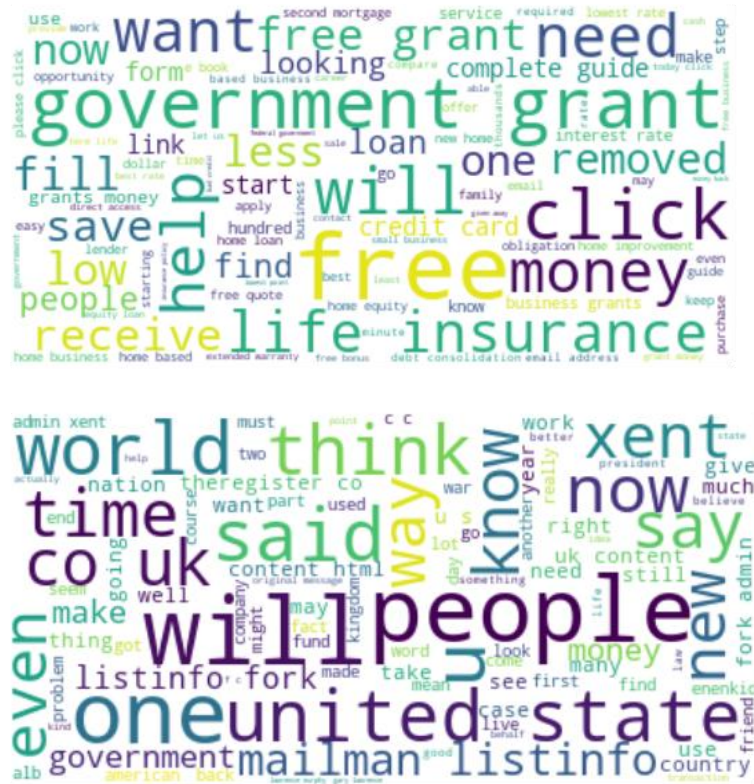


Figure 3-Sample Word Cloud of Email Texts

Naïve Bayes Classification

It is a classification technique based on Bayes' Theorem with an assumption of independence among predictors. In simple terms, a Naïve Bayes classifier assumes that the presence of a particular feature in a class is unrelated to the presence of any other feature. This study uses NB Classification to classify spam emails.

3. Results

Naïve-Bayes algorithm is run with Grid search with several parameters, to find the optimal results.

Mean	Std-dev	Alpha	Fit Prior
0.985567	-0.002797	0.0001	TRUE
0.986957	-0.002473	0.0001	FALSE
0.98717	-0.002339	0.001	TRUE
0.988774	-0.002253	0.001	FALSE
0.987064	-0.002159	0.01	TRUE
0.989202	-0.002763	0.01	FALSE
0.980862	-0.00194	0.1	TRUE
0.983642	-0.002909	0.1	FALSE
0.837806	-0.008439	1	TRUE
0.882925	-0.007082	1	FALSE
0.743505	-0.000318	10	TRUE
0.747675	-0.003525	10	FALSE
0.743505	-0.000318	100	TRUE
0.744039	-0.000549	100	FALSE
0.743505	-0.000318	1000	TRUE

Table 3-NB Classifier Grid Search Results (Green highest accuracy)

Best Accuracy with Grid Search: 0.989

Training data Metrics

Metrics	Value
average accuracy	0.993
average precision	0.997
average recall	0.977

Test data Metrics

Metrics	Value
average accuracy	0.989
average precision	0.990
average recall	0.968

Classification report

	Precision	Recall	F1-Score	Support
1 (Spam)	0.99	0.97	0.98	2399
0 (Not Spam)	0.99	1.00	0.99	6954
Accuracy			0.99	9353
Macro avg	0.99	0.98	0.99	9353
Weighted avg	0.99	0.99	0.99	9353

ROC Curve




Figure 4:ROC Curve (True Positive vs False Positive)

The ROC curve for the NB classifier showed a mean area under the curve AUC of 1.00.

The accuracy, precision, and f-score are all very high and close to 1.00. The ROC curve also indicates there are no false positives compared to the true positive rate. The area under the curve indicates the results cover the entire dataset.

Conclusion

The features (texts) indicate that spam vs non-spam is: The positive feature_coef with high values corresponds to a higher incidence of spams vs feature_coef (Negative) with lowest scores corresponding to least incidence of spams.

features	feature_coef (Positive)	features	feature_coef (Negative)
sightings	8.086723	date	-10.818954
please	7.802528	com	-9.38241
click	6.755935	newsisfree	-8.952799
remove	6.451776	wrote	-8.711082

free	6.294921	rpm	-7.929359
money	6.189775	supplied	-6.946106
removed	5.814065	said	-6.220962
visit	4.46381	cnet	-5.064446
receive	4.329541	use	-4.781836
offer	4.261584	exmh	4.603416

Figure 5-Feature Importance based on co-efficient

The above table displays the features with their coefficient levels. These are the top 10 (positive & negative) coefficients after the L2 penalty. The table shows the common words in spam (money, free, offer, click) do tend to appear more on spam messages. On the contrary, the messages with negative coef words like cnet, use, date do tend to appear more on non-spam messages.

This model is recommended to use as it produces higher accuracy and precision (with a high f-score).

Appendix – Code

Nbviewer link:

<https://nbviewer.org/github/ravisiv/SpamClassificationML/blob/main/Case%20Study%203%20Balaji%20-%20Apurv%20-%20Ravi.ipynb>

Python Jupyter Notebook

Case Study 3: Spam classifier

Submitted by:

- Ravi Sivaraman
- Balaji Avvaru
- Apurv Mittal

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import os
from os.path import isfile
import email
#import BeautifulSoup
from bs4 import BeautifulSoup
import re
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import silhouette_score
import hdbscan
from sklearn.cluster import KMeans
from sklearn.model_selection import cross_val_predict
```

```

from sklearn.metrics import classification_report
from sklearn.model_selection import cross_validate
from sklearn.linear_model import LogisticRegression
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics import accuracy_score
from sklearn.model_selection import StratifiedKFold
from sklearn.model_selection import GridSearchCV
import seaborn as sns
from sklearn.preprocessing import label_binarize
from sklearn import metrics as mt
from sklearn.feature_extraction import text
from wordcloud import WordCloud
from scipy.sparse import hstack

import warnings
warnings.filterwarnings("ignore")
# location of emails
data_path = "/Users/ravis/Downloads/SpamAssassinMessages"
# get all sub folders
sub_folders = [x[0] for x in os.walk(data_path) if x[0] != data_path]
%%time

# read all emails from all sub folders
mail_ty = []
text_ty = []
data = []
target = []
email_attachment = []
attachment = False

for folder in sub_folders:
    files = [f for f in os.listdir(folder) if os.path.isfile(os.path.join(folder, f))]
    for file in files:
        with open(f"{folder}/{file}", encoding="latin1") as f:
            # with open(f"{folder}/{file}", "r") as f:
            x = email.message_from_file(f)
            # print(x)
            # if (file != 'cmds'):
            #     mail_data.append(lines)

            mail_type = x.get_content_type()
            text_type = x.get_content_charset()
            mail_ty.append(mail_type)
            text_ty.append(text_type)
            if re.search("spam", folder):
                target.append(1)
            else:
                target.append(0)

            if mail_type == "text/html":
                if not (isinstance(x.get_payload(), str)) and x.get_payload().get('Content-
Disposition'):
                    dispositions = x.get_payload().get("Content-Disposition",
None).strip().split(";")
                    if bool(dispositions[0].lower() == "attachment"):
                        attachment = True
                    else:
                        attachment = False

                    tmp = BeautifulSoup(x.get_payload(), 'html.parser')
                    tmp = tmp.text.replace("\n", " ")
                    data.append(tmp)
            elif "multipart" in mail_type:
                attachment = False
                multipart_data = []
                for text in x.get_payload():
                    if not isinstance(text, str):
                        if text.get('Content-Disposition'):
                            dispositions = text.get("Content-Disposition", None).strip().split(";")
                            if bool(dispositions[0].lower() == "attachment"):
                                attachment = True

```

```

        if text.get_content_type() == "text/html":
            tmp = BeautifulSoup(text.get_payload(), 'html.parser')
            tmp = tmp.text.replace("\n", " ")
            multipart_data.append(tmp)
        elif text.get_content_type() == "text/plain":
            multipart_data.append(text.get_payload())

    multipart_email = [''.join(str(item)) for item in multipart_data]
    data.append(multipart_email)
else:
    if not (isinstance(x.get_payload(), str)) and x.get_payload().get('Content-
Disposition'):
        dispositions = x.get_payload().get("Content-Disposition",
None).strip().split(";")
        if bool(dispositions[0].lower() == "attachment"):
            attachment = True
        else:
            attachment = False
        data.append(x.get_payload())

    if attachment:
        email_attachment.append(1)
    else:
        email_attachment.append(0)

# Reference: https://gaurav.kuwar.us/index.php/2017/10/09/extracting-files-from-raw-email-with-
python/
df = pd.DataFrame()
df["mail_types"] = mail_ty
df["text_types"] = text_ty
# Count of mail types
df["mail_types"].value_counts()
# Count of text types
df["text_types"].value_counts()
# Create a data frame with email text and target (whether mail is spam or not, 1 for spam and 0
for not a spam)
email_df = pd.DataFrame()
email_df["data"] = data
#email_df["mail_type"] = mail_ty
#email_df["text_type"] = text_ty
email_df["target"] = target
email_df["Attachments"] = email_attachment
email_df['target'].value_counts()
email_df.info()
email_df.loc[1].data
email_df["data_new"] = [''.join(str(item).lower()) for item in email_df.data]
print(email_df["data_new"][0])
# get the instanc of TfidfVectorizer
#my_stop_words = text.ENGLISH_STOP_WORDS.union(["spamassassin", "email", "message", "\n", "nbsp",
"font", "exhm", "subject", "list", "url", "net"])
from nltk.corpus import stopwords
stop = list(stopwords.words('english'))
stop.extend("spamassassin email message \n nbsp font exhm subject list url net http www org html
linux 2002 font e2 c2 div 0d c2 0a xa0 8c 2ffont e2 3e sourceforge spamassasin 01 yahoo 1440
a0".split())

tf_vectorizer = TfidfVectorizer(analyzer = 'word', stop_words=set(stop))

# tf_vectorizer = TfidfVectorizer(ngram_range=(1,2), stop_words=text.ENGLISH_STOP_WORDS)

#tf_vectorizer = TfidfVectorizer()
# fit and transform email data
new_vectors = tf_vectorizer.fit_transform(email_df.data_new)
# Pie chart
plt.figure(figsize=(5,4))
email_df.target.value_counts().plot.pie(autopct = "%.1f%%")
plt.title("Proportion of Target Value")
plt.show()
email_df['Attachments'].value_counts()
plt.figure(figsize=(5,4))

```

```
sns.countplot(x="Attachments", data = email_df)
plt.title("Distribution of Attachments")
plt.show()
plt.figure(figsize=(5,4))
sns.countplot(x="Attachments", hue = "target", data = email_df)
plt.title("Attachments in Spam (1) vs Not Spam (0)")
plt.show()
new_vectors = hstack((new_vectors,np.array(email_attachment)[: ,None]))
```

Clustering

KMeans Clustering

K-means clustering is a type of unsupervised learning, which is used when you have unlabeled data (i.e., data without defined categories or groups). The goal of this algorithm is to find groups in the data, with the number of groups represented by the variable K. The algorithm works iteratively to assign each data point to one of K groups based on the features that are provided. Data points are clustered based on feature similarity

KMeans Clustering with default parameters

```
wcss = []
score = []
K = range(2,30)
for k in K:
    km = KMeans(n_clusters=k, random_state=1234, init = 'k-means++')
    km = km.fit(new_vectors)
    labels = km.predict(new_vectors)
    wcss.append(km.inertia_)
    sc = silhouette_score(new_vectors, labels)
    score.append(sc)

plt.rcParams['figure.figsize'] = (15, 5)
#plt.subplot(1,1,1)
plt.plot(K, wcss, 'bx-')
plt.xlabel('Number of centroids')
plt.ylabel('Within-Cluster-Sum-of-Squares')
plt.title('Elbow Method For Optimal k')
```

Visualize Clusters

We will use a technique called t-SNE (t-distributed Stochastic Neighbor Embedding) to generate a 2 dimensional representation of our dataset, in order to have a more intuitive understanding of how the clustering looks. First let's look at an un-clustered version of this 2D projection.

```
%%time
from sklearn.manifold import TSNE
import sklearn.cluster as cluster
newdims = (12, 8)
plt.subplots(1, 1, figsize=newdims)
plt.subplot(1, 1, 1)
plot_kwds = {'alpha' : 0.25, 's' : 40, 'linewidths':0}
projection = TSNE().fit_transform(new_vectors)
plt.scatter(*projection.T, **plot_kwds)
plt.title("")
plt.show()
```

Now look at clustered version of this 2D projection with various clustering techniques

```
%%time
import seaborn as sns

# This function will run a given clustering algorithm and plot the clusters on the same 2D TSNE
projection as above
def plot_clusters(data, algorithm, args, kwds):
    labels = algorithm(*args, **kwds).fit_predict(data)
    palette = sns.color_palette('muted', np.unique(labels).max() + 1)
```

```

colors = [palette[x] if x >= 0 else (.5, .5, .5) for x in labels]
plt.scatter(*projection.T, s=50, linewidth=0, c=colors, alpha=0.25)
frame = plt.gca()
frame.axes.get_xaxis().set_visible(True)
frame.axes.get_yaxis().set_visible(True)
#plot_kws = {'alpha' : 0.25, 's' : 40, 'linewidths':0}
plt.title('Clusters found by {}'.format(str(algorithm.__name__)), fontsize=14)

best_k = 15
# plot the clusters
newdims = (10, 8)
plt.subplots(1, 1, figsize=newdims)
plt.subplot(1, 1, 1)
plot_clusters(new_vectors, cluster.KMeans, (), {'n_clusters':best_k})

```

Word Cloud

```

# word cloud with best K
km = KMeans(n_clusters=best_k, init = 'k-means++')
km = km.fit(new_vectors)
labels = km.predict(new_vectors)

clusters = list(labels)

kmeans_result={'cluster':clusters,'reviews':email_df.data_new}
kmeans_result=pd.DataFrame(kmeans_result)
for k in range(0,12):
    s=kmeans_result[kmeans_result.cluster==k]
    text=s['reviews'].str.cat(sep=' ')
    text=text.lower()
    text=' '.join([word for word in text.split()])
    wordcloud = WordCloud(max_font_size=50, max_words=100,
background_color="white").generate(text)
    plt.figure()
    plt.imshow(wordcloud, interpolation="bilinear")
    plt.axis("off")
    plt.show()

```

Naive Bayes Classification

It is a classification technique based on Bayes' Theorem with an assumption of independence among predictors. In simple terms, a Naive Bayes classifier assumes that the presence of a particular feature in a class is unrelated to the presence of any other feature.

```

def displayModel_metrics(best_model, grid_model, features, target, cv):
    metrics = cross_validate(best_model, features, y=target, cv=cv,
                             scoring=['accuracy','precision','recall'], return_train_score=True)

    y_predict = cross_val_predict(best_model, features, target, cv=cv)

    print('\nBest Accuracy with Grid Search           : {:.3f}'.format(grid_model.best_score_))
    print('\nTraining data Metrics')
    print('\n    The average accuracy : {:.3f}'.format(metrics['train_accuracy'].mean()))
    print('    The average precision : {:.3f}'.format(metrics['train_precision'].mean()))
    print('    The average recall    : {:.3f}'.format(metrics['train_recall'].mean()))

    print('\nTest data Metrics')
    print('\n    The average accuracy : {:.3f}'.format(metrics['test_accuracy'].mean()))
    print('    The average precision : {:.3f}'.format(metrics['test_precision'].mean()))
    print('    The average recall    : {:.3f}'.format(metrics['test_recall'].mean()))

    matrix = classification_report(target, y_predict, labels=[1,0])
    print('\nClassification report\n')
    print(matrix)

# Reference
https://github.com/jakemdw/DataMiningNotebooks/blob/master/06.%20Classification.ipynb
# ROC curve plot

```

```

def roc_curve_plot(model_fit, features, target):

    sns.set_palette("dark")

    yhat_score = model_fit.predict_proba(features)

    # Compute ROC curve for a subset of interesting classes
    fpr = dict()
    tpr = dict()
    roc_auc = dict()
    for i in np.unique(target):
        fpr[i], tpr[i], _ = mt.roc_curve(target, yhat_score[:, i], pos_label=i)
        roc_auc[i] = mt.auc(fpr[i], tpr[i])

    for i in np.unique(target):
        plt.plot(fpr[i], tpr[i], label= ('class %d (area = %0.2f)' % (i, roc_auc[i])))
    plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')

    plt.legend(loc="lower right")
    plt.title('Receiver operating characteristic')
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.show()

#Create Cross Validation Procedure
cv = StratifiedKFold(n_splits=10, random_state=1234, shuffle=True)
# Naive Bayes (NB) classifier
clf = MultinomialNB().fit(new_vectors,email_df['target'])
# define parameters
C_nb = [0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000]
nb_prior=[True, False]

nb_clf = MultinomialNB()
# define grid search
param_grid_nb = dict(alpha=C_nb, fit_prior=nb_prior)

grid_search_nb = GridSearchCV(estimator=nb_clf, param_grid=param_grid_nb, n_jobs=-1, cv=cv,
                              scoring='accuracy',error_score=0)

%%time
grid_result_nb = grid_search_nb.fit(new_vectors,email_df['target'])
# summarize results
print("Best: %f using %s" % (grid_result_nb.best_score_, grid_result_nb.best_params_))
means = grid_result_nb.cv_results_['mean_test_score']
stds = grid_result_nb.cv_results_['std_test_score']
params = grid_result_nb.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))
# The GridSearch algorithm determined the following optimal parameters
best_Estimator_nb =grid_result_nb.best_estimator_
best_Estimator_nb
# Display model metrics
displayModel_metrics(best_Estimator_nb, grid_result_nb, new_vectors,email_df['target'], cv)
# Plot ROC curve
roc_curve_plot(grid_result_nb, new_vectors, email_df['target'])

```

Naive Bayes Classification with clusters as feature

```

# add clusters as feature
new_vectors = hstack((new_vectors,np.array(clusters)[: ,None]))
new_vectors
%%time
grid_result_nb = grid_search_nb.fit(new_vectors,email_df['target'])
# summarize results
print("Best: %f using %s" % (grid_result_nb.best_score_, grid_result_nb.best_params_))
means = grid_result_nb.cv_results_['mean_test_score']
stds = grid_result_nb.cv_results_['std_test_score']
params = grid_result_nb.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))
# The GridSearch algorithm determined the following optimal parameters
best_Estimator_nb =grid_result_nb.best_estimator_
best_Estimator_nb

```

```
# Display model metrics
displayModel_metrics(best_Estimator_nb, grid_result_nb, new_vectors,email_df['target'], cv)
# Plot ROC curve
roc_curve_plot(grid_result_nb, new_vectors, email_df['target'])
```

Feature importance with Logistic regression

```
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
LR = LogisticRegression()

# define parameters
penalty_LR = ['l1', 'l2', 'elasticnet', 'none']
#penalty_LR = [ 'l1', 'l2']
C_LR = [0.001, 0.01, 0.1, 1, 10, 100, 1000]
#C_LR = [0.001,10, 100]
max_iter_LR = [500]
#max_iter_LR = [500]
class_weight_LR = ['balanced']
#solver_LR = ['newton-cg', 'lbfgs', 'liblinear', 'sag', 'saga']
solver_LR = ['lbfgs', 'liblinear']

# define grid search
param_grid_LR = dict(penalty=penalty_LR, C=C_LR, max_iter=max_iter_LR,
class_weight=class_weight_LR, solver=solver_LR)

grid_search_LR = GridSearchCV(estimator=LR, param_grid=param_grid_LR, n_jobs=-1, cv=cv,
scoring='accuracy',error_score=0)

%%time
grid_result_LR = grid_search_LR.fit(new_vectors,email_df['target'])
# summarize results
print("Best: %f using %s" % (grid_result_LR.best_score_, grid_result_LR.best_params_))
means = grid_result_LR.cv_results_['mean_test_score']
stds = grid_result_LR.cv_results_['std_test_score']
params = grid_result_LR.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))
# The GridSearch algorithm determined the following optimal parameters
best_Estimator_LR =grid_result_LR.best_estimator_
best_Estimator_LR
features = tf_vectorizer.get_feature_names()
features.append('email_attachment')
features.append('clusters')

feature_importance_df = pd.DataFrame(features, columns=['features'])
feature_importance_df['feature_coef'] = best_Estimator_LR.coef_[0]

feature_importance_df.head()
feature_importance_df = feature_importance_df.sort_values(by=['feature_coef'])
feature_importance_df.tail(20)
feature_importance_df.head(20)
```