



Case study 6

Artificial Neural Networks

Balaji Avvaru, Apurv Mittal, Ravi
Sivaraman

Quantifying The World

Introduction

The goal of this project is to maximize accuracy and minimize loss using a dense neural network. The dataset contains 7 million records (big dataset), with 29 features. The features are just numbered from f0 to f26, label. “# label” is the target variable in the dataset.

Data Analysis

The dataset contains 7 million records (big dataset), with 29 features. The features are numbered from f_0 to f_{26} , label. “# label” is the target variable in the dataset. There are no missing values. This dataset doesn’t require imputation for the features.

The correlation matrix below shows the collinearity among the variables. Data shows there is high collinearity for a few of the attributes.

This study considered all the features with lower than 99% correlation for this analysis. There are features with up to 83% of correlation. Since none of the features showed 99% or above collinearity, all features are considered for modeling, and none are dropped.

Correlation Analysis




Figure 1: Collinearity heatmap of Attributes

Target

Target is a binary variable called '*# label*', which has value either *0* or *1*.




Figure 2: Target variable distribution plot

2. Methods

f9, f13, f17, f21 are categorical variables, which are converted using one-hot encoding.

All the features are normalized using the `StandardScaler` technique. After normalization, the data was split into train and test datasets by a factor of 80/20.

The training data was further split into training and validation set by 80/20. The validation set is used to validate the model performance during training.

The test data is unseen data and is used to calculate the final accuracy of the best model and ROC curve, and confusion matrix.

ANN

Artificial neural networks (ANNs) are models inspired by humans (or biological) neural networks.

An ANN is a collection of neurons (nodes) loosely based on the biological brain interconnected in a dense layer. Each neural is a linear regression created from each feature. The neurons send a signal once it crosses a certain threshold. Usually, they are aggregated into several layers (sequential or dense). The final layer depends on the type of classification required. For binary classification, sigmoid functions are used. The last layer uses a sigmoid function for this study, as the target is a binary classification.

Layers of Neural Network

In this study, the authors created two sets of neural networks, one with 3 layers and another with 4 layers.

First neural network with the following layers:

Layer	Activation Function	Number of Neurons	Input to Layer
1	gelu	512	Features
2	gelu	256	Layer1
3	gelu	128	Layer2
4	sigmoid	2	Layer3

Table 1: Layers of first Neural Network

Second neural network with the following layers:

Layer	Activation Function	Number of Neurons	Input to Layer
1	gelu	1024	Features
2	gelu	512	Layer1
3	gelu	256	Layer2
4	gelu	128	Layer3
5	sigmoid	2	Layer4

Table 2: Layers of second Neural Network

The hyper-parameters (tunable parameters) are:

Third neural network with the following layers:

Layer	Activation Function	Number of Neurons	Input to Layer
1	swish	512	Features
2	swish	256	Layer1
3	swish	128	Layer2
4	sigmoid	2	Layer3

Table 3: Layers of third Neural Network

Parameters:

batchsize: Constant that multiplies the regularization term; the higher the value, the stronger the regularization. Also used to compute the learning rate when *learning_rate* is set to 'optimal.'

Optimizer: Adam(lr=1e-2).

loss: Sparse Categorical Cross-Entropy

safety: Early Stopping (using *val_loss*)

patience: 3

min_delta: 2e-4

The models are compiled to run for 50 epochs. Each epoch runs the entire dataset once. With 50 epochs, the model runs each record in the dataset 50 times or unless early stopping condition is met (when loss stops reducing further). If the accuracy or *val_loss* doesn't improve, then the model doesn't run the entire 50 epochs and stops early. The model is configured to stop if 3 (hyper-param *patience*) consecutive epochs did not improve *val_loss*. This prevents running neural networks when the metrics don't improve.

3. Results

Model	Batch Size	Dense Neurons	Validation Accuracy
gelu/4 Layers	2048	1094 → 512 → 128	87.99%
gelu/3 Layers	2048	512 → 256	88.20%
gelu/3 Layers	1024	512 → 256	88.07%
swish/3 Layers	2048	512 → 256	88.12%
gelu/3 Layers	100	512 → 256	87.63%

Table 4: Results of various models

All the models are giving similar results. The *gelu* model with 3 layers and batch size of 2048 gives the highest accuracy of 88.20%.

When batch size decreases, training time increases (when all others are kept same):

Batch Size	Training Time (in seconds)
2048	16
1024	31
100	250

Table 5: Training Times by Batch Size

The detailed results from the best of the models (*gelu/3layers/2048 batch*) are listed below:

	precision	recall	f1-score	support
0 (target)	0.9	0.85	0.88	699445
1 (target)	0.86	0.91	0.88	700555
accuracy			0.88	1400000
macro avg	0.88	0.88	0.88	1400000
weighted avg	0.88	0.88	0.88	1400000

Table 6: Confusion Matrix




Figure 3: Confusion Matrix

4. Conclusion

All neural network models produced similar results. In terms of accuracy, there is no significant difference in the model output. Changing the activation function between `gelu` or `swish` doesn't change the final accuracy significantly. Similarly, the addition of the 4th layer doesn't produce higher accuracy results. The change in batch size between 1024 and 2048 also doesn't produce a significantly different output.

Due to early stopping criteria, the model stops much earlier than running all 50 epochs as it stops improving further. Overall model accuracy is around 88.20%.

The neural network model was able to accurately predict the existence of new particles with 88% accuracy, we can advise our clients to use neural network model to detect a new particle.

Appendix – Code

NB Viewer Link:

https://nbviewer.org/github/ravisiv/CaseStudy6_NN/blob/main/Case%20Study%206.ipynb

CaseStudy6_NN (/github/ravisiv/CaseStudy6_NN/tree/main)

/ Case Study 6.ipynb (/github/ravisiv/CaseStudy6_NN/tree/main/Case Study 6.ipynb)

```
In [1]: import os
import
pandas as pd
import re
import
datetime as
dt import
numpy as np
from IPython.display

import display import
warnings import seaborn
as sns
from sklearn.preprocessing import
StandardScaler from sklearn import metrics
as mt
from sklearn.metrics import
classification_report from sklearn.metrics
import f1_score import
matplotlib.pyplot as plt
from sklearn.linear_model import
SGDClassifier
from sklearn.metrics import precision_recall_curve,
plot_precision_reca from sklearn.preprocessing import
label_binarize, StandardScaler from sklearn import metrics as mt
from sklearn.model_selection import cross_validate,
cross_val_predict, warnings.filterwarnings('ignore')
```

```
In [2]: # load data
#df = pd.read_csv('all_train.csv') # read in the csv file
df = pd.read_csv('/Users/ravis/Downloads/data/all_train.csv')
```

```
In [3]: df.mass
```

```
Out[3]: 0    1000.000000  1
        750.000000
        2          750.000000
        3    1250.000000
        4          750.000000  ...
```



```

6999995      750.000000
6999996     1250.000000
6999997     1500.000000
6999998     1500.000000
6999999      499.999969
Name: mass, Length: 7000000, dtype: float64

```

In [4]:

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'> RangeIndex:
7000000 entries, 0 to 6999999
Data columns (total 29 columns):
#   Column  Dtype
---  -
0    # label  float64
1    f0       float64
2    f1       float64
3    f2       float64
4    f3       float64
5    f4       float64
6    f5       float64
7    f6       float64
8    f7       float64
9    f8       float64
10   f9       float64
11   f10      float64
12   f11      float64
13   f12      float64
14   f13      float64
15   f14      float64
16   f15      float64
17   f16      float64
18   f17      float64
19   f18      float64
20   f19      float64
21   f20      float64
22   f21      float64
23   f22      float64
24   f23      float64
25   f24      float64
26   f25      float64
27   f26      float64 28 mass    float64 dtypes: float64(29) memory usage: 1.5 GB In [5]:

```

```
df.describe()
```

Out[5]:

	# label	f0	f1	f2	f3	f4
count	7.000000e+06	7.000000e+06	7.000000e+06	7.000000e+06	7.000000e+06	7.000000e+06
mean	5.001256e-01	1.612528e-02	4.770022e-04	2.686578e-05	1.056081e-02	-1.050026e-04
std	5.000000e-01	1.004417e+00	9.974864e-01	1.000080e+00	9.956003e-01	9.998670e-01
min	0.000000e+00	-1.960549e+00	-2.365355e+00	-1.732165e+00	-9.980274e+00	-1.732137e+00
25%	0.000000e+00	-7.288206e-01	-7.332548e-01	-8.656704e-01	-6.092291e-01	-8.658025e-01
50%	1.000000e+00	-3.930319e-02	8.523957e-04	3.199154e-04	1.963316e-02	-5.070131e-04
75%	1.000000e+00	6.900799e-01	7.347832e-01	8.659464e-01	6.798818e-01	8.657646e-01
max	1.000000e+00	4.378282e+00	2.365287e+00	1.732370e+00	4.148023e+00	1.731978e+00

8 rows × 29 columns

Missing value analysis

In [6]:

Out[6]: In 0
[7]: **Target**

Out[7]: 1.0 3500879
0.0 3499121
Name: # label, dtype: int64

```
df[ '# label' ].value_counts ()
```

```
# Validate null values in the csv file
df.isnull () .sum() .sum()
```


In [8]:

```
sns . countplot ( x = "# label" , data = df )  
plt . title ( "Distribution of Target Values" )  
plt . show()
```



In [9]:

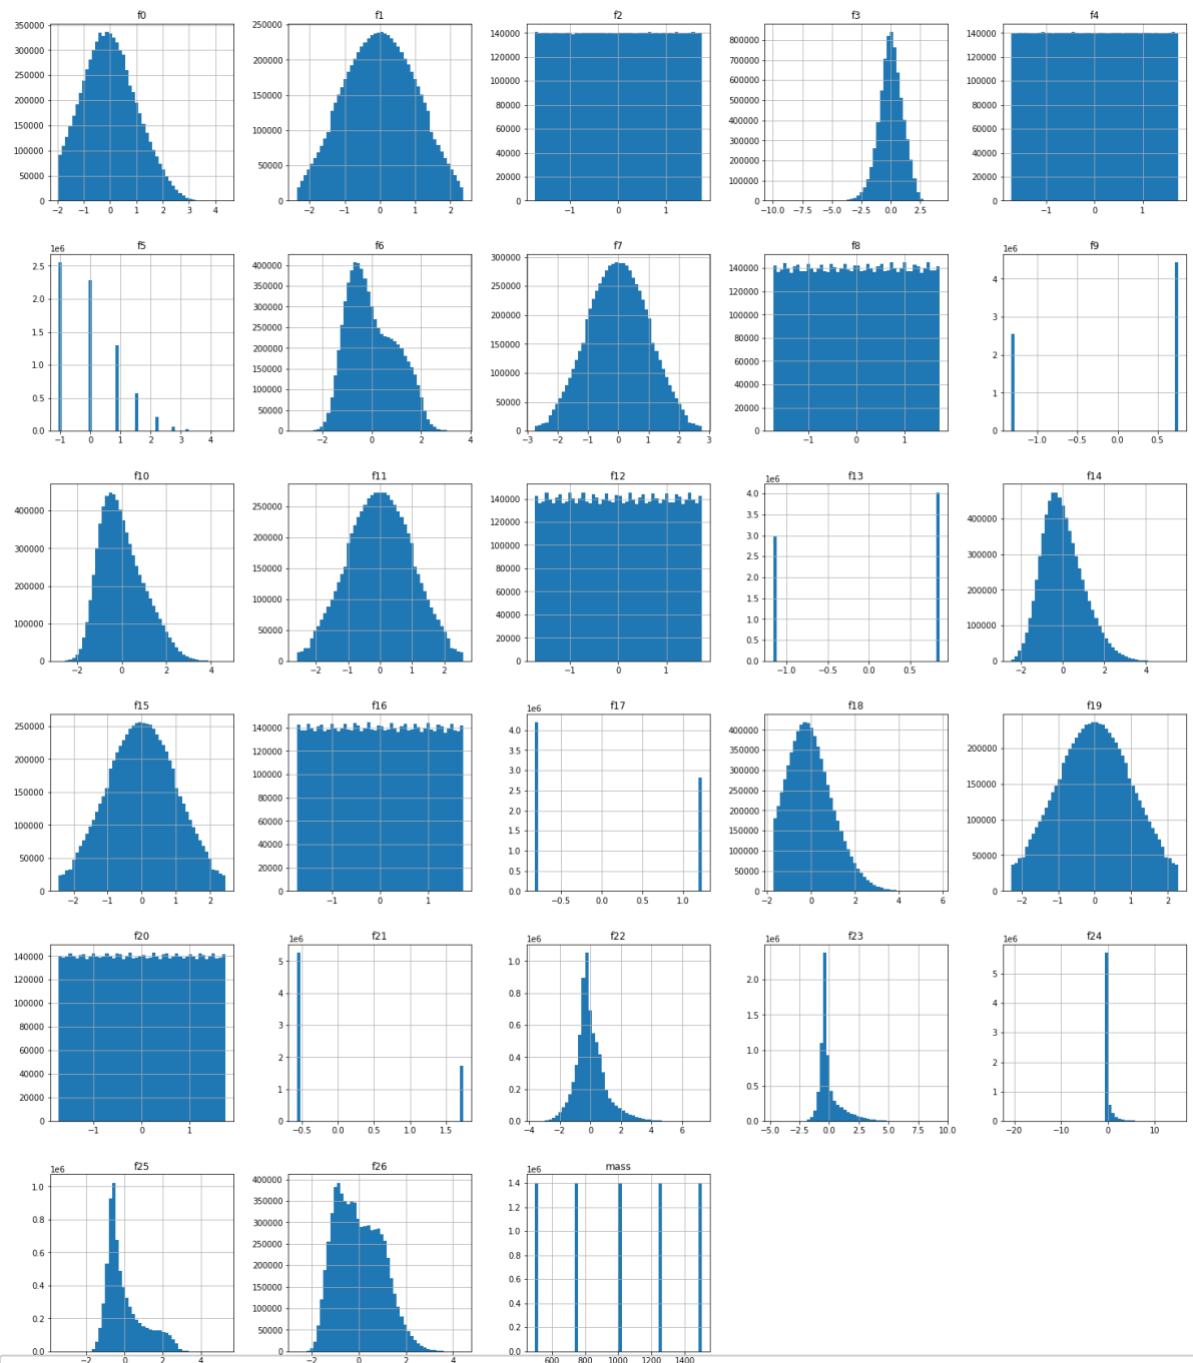
```
# Pie chart
df['# label'].value_counts().plot.pie(autopct = "%1f%%")
plt.title("Proportion of Target Value")
plt.show()
```



Independent Variable analysis p y

In [10]:

```
#Visualizing the hist of data to check normality of independent v
df_X = df.drop(['# label' ], axis =1)
df_X.hist ( bins =50, figsize =( 25, 30))
plt . show()
```



In [11]:

```
print(df['f9'].value_counts()) print(df['f13'].value_counts())
print(df['f17'].value_counts()) print(df['f21'].value_counts())
```

```

0.754261  4438579 -1.325801  2561421
Name: f9, dtype: int64
0.860649  4027351 -1.161915  2972649
Name: f13, dtype: int64
-0.815440  4187343
1.226331  2812657 Name: f17, dtype: int64
-0.573682  5265796
1.743123  1734204 Name: f21, dtype: int64

```

```
df.drop(['# label' ], axis =1).describe ()
```

In [12]:


Out[12]:

	f0	f1	f2	f3	f4	f5
count	7.000000e+06	7.000000e+06	7.000000e+06	7.000000e+06	7.000000e+06	7.000000e+06
mean	1.612528e-02	4.770022e-04	2.686578e-05	1.056081e-02	-1.050026e-04	2.765919e-0
std	1.004417e+00	9.974864e-01	1.000080e+00	9.956003e-01	9.998670e-01	1.000957e+0
min	-1.960549e+00	-2.365355e+00	-1.732165e+00	-9.980274e+00	-1.732137e+00	-1.054221e+0
25%	-7.288206e-01	-7.332548e-01	-8.656704e-01	-6.092291e-01	-8.658025e-01	-1.054221e+0
50%	-3.930319e-02	8.523957e-04	3.199154e-04	1.963316e-02	-5.070131e-04	-5.983562e-03
75%	6.900799e-01	7.347832e-01	8.659464e-01	6.798818e-01	8.657646e-01	8.504885e-0
max	4.378282e+00	2.365287e+00	1.732370e+00	4.148023e+00	1.731978e+00	4.482618e+0

8 rows x 28 columns

In [13]:

```
#heatmap - correlation matrix
plt .figure (figsize =( 60, 60)) #code reference (5-1)
plt .xticks (rotation =90, fontsize = 35)
plt .yticks (rotation =180, fontsize = 35)
ax=sns . heatmap( df_X . corr () , annot =False, cbar =True, annot_kws={"size" : 2}
cbar = ax . collections [ 0] . colorbar
cbar . ax . tick_params (labelsize =30)
plt . title ( 'HeatMap-Correlation Matrix' , fontsize = 45)
plt . show()
```



Check for Multicolliniarity

```
In [14]: #https://www.projectpro.io/recipes/drop-out-highly-  
correlated-features-  
# to drop features with colliniarity more than 95%  
pd.set_option('display.max_rows', 100)
```

```
corr_df = pd.DataFrame(df_X.corr().abs()) corr_df.head(100)
```

Out[14]:

	f0	f1	f2	f3	f4	f5	f6		
	f7	f							
f0	1.000000	0.000556	0.000321	0.012037	0.000464	0.078401	0.349973	0.000026	0.00092
f1	0.000556	1.000000	0.000200	0.000706	0.000131	0.000511	0.000454	0.315357	0.00058
f2	0.000321	0.000200	1.000000	0.000074	0.174967	0.000162	0.000436	0.000024	0.23231
f3	0.012037	0.000706	0.000074	1.000000	0.000385	0.092129	0.468157	0.000091	0.00018
f4	0.000464	0.000131	0.174967	0.000385	1.000000	0.000496	0.000307	0.000435	0.23300
f5	0.078401	0.000511	0.000162	0.092129	0.000496	1.000000	0.191900	0.000740	0.00067
f6	0.349973	0.000454	0.000436	0.468157	0.000307	0.191900	1.000000	0.000491	0.00069
f7	0.000026	0.315357	0.000024	0.000091	0.000435	0.000740	0.000491	1.000000	0.00088
f8	0.000924	0.000580	0.232319	0.000188	0.233003	0.000671	0.000699	0.000887	1.00000
f9	0.039924	0.000141	0.000426	0.021872	0.000084	0.143509	0.038961	0.000328	0.00054
f10	0.171641	0.000726	0.000291	0.306169	0.000229	0.243807	0.799686	0.000724	0.00048
f11	0.000564	0.336327	0.000268	0.000073	0.000174	0.000271	0.000290	0.306276	0.00046
f12	0.000268	0.000218	0.036190	0.000460	0.001119	0.000454	0.000685	0.000132	0.23824
f13	0.057995	0.000088	0.000378	0.024117	0.000092	0.137170	0.000264	0.000083	0.00012
f14	0.147686	0.000375	0.000017	0.244698	0.000189	0.363165	0.575503	0.000471	0.00040
f15	0.000105	0.256024	0.000185	0.001113	0.000011	0.000005	0.000920	0.324219	0.00046
f16	0.000684	0.000249	0.091000	0.000181	0.065579	0.000300	0.000205	0.000238	0.10618
f17	0.027194	0.000136	0.000912	0.008179	0.000023	0.042581	0.038356	0.000435	0.00047
f18	0.119041	0.000079	0.000026	0.174068	0.000338	0.517752	0.390014	0.000251	0.00023
f19	0.000152	0.195312	0.000399	0.000328	0.000400	0.000337	0.000226	0.299012	0.00012
f20	0.000617	0.000254	0.092043	0.000047	0.068271	0.000305	0.000303	0.000276	0.05167
f21	0.034794	0.000191	0.000312	0.010787	0.000485	0.044609	0.034123	0.000131	0.00004
f22	0.080169	0.000100	0.000171	0.096627	0.000136	0.195117	0.209496	0.000023	0.00009
f23	0.109947	0.000384	0.000075	0.190014	0.000061	0.188721	0.433089	0.000133	0.00035

f24	0.289670	0.000178	0.000499	0.056976	0.000009	0.050882	0.126159	0.000182	0.00065
f25	0.296976	0.000563	0.000186	0.406295	0.000346	0.082941	0.581645	0.000020	0.00037
f26	0.327533	0.000782	0.000315	0.482503	0.000414	0.308311	0.835995	0.000621	0.00029
mass	0.126717	0.000050	0.000169	0.203470	0.000002	0.064497	0.310024	0.000069	0.00075

28 rows × 28 columns



In [15]:

```
# Multi Colliniarity analysis on Independent variables
upper_tri = corr_df.where(np.triu(np.ones(corr_df.shape),k=1).astype(np
print(upper_tri)
```

```
f0    f1    f2    f3    f4    f5    f6 \ f0  NaN  0.000556  0.000321  0.012037
0.000464  0.078401  0.349973  f1  NaN    NaN  0.000200  0.000706  0.000131  0.000511
0.000454  f2  NaN    NaN    NaN  0.000074  0.174967  0.000162  0.000436  f3  NaN    NaN
NaN    NaN  0.000385  0.092129  0.468157  f4  NaN    NaN    NaN    NaN    NaN
0.000496  0.000307  f5  NaN    NaN    NaN    NaN    NaN    NaN  0.191900  f6  NaN
NaN    NaN    NaN    NaN    NaN    NaN  f7  NaN    NaN    NaN    NaN    NaN
NaN    NaN  f8  NaN    NaN    NaN    NaN    NaN    NaN    NaN  f9  NaN    NaN
NaN    NaN    NaN    NaN    NaN  f10  NaN    NaN    NaN    NaN    NaN    NaN
NaN  f11  NaN    NaN    NaN    NaN    NaN    NaN    NaN  f12  NaN    NaN    NaN
NaN    NaN    NaN    NaN  f13  NaN    NaN    NaN    NaN    NaN    NaN  f14
NaN    NaN    NaN    NaN    NaN    NaN    NaN  f15  NaN    NaN    NaN    NaN
NaN    NaN    NaN  f16  NaN    NaN    NaN    NaN    NaN    NaN    NaN  f17  NaN
NaN    NaN    NaN    NaN    NaN    NaN  f18  NaN    NaN    NaN    NaN    NaN
NaN    NaN  f19  NaN    NaN    NaN    NaN    NaN    NaN    NaN  f20  NaN    NaN
NaN    NaN    NaN    NaN    NaN  f21  NaN    NaN    NaN    NaN    NaN    NaN
NaN  f22  NaN    NaN    NaN    NaN    NaN    NaN    NaN  f23  NaN    NaN    NaN
NaN    NaN    NaN    NaN  f24  NaN    NaN    NaN    NaN    NaN    NaN  f25
NaN    NaN    NaN    NaN    NaN    NaN    NaN  f26  NaN    NaN    NaN    NaN
NaN    NaN    NaN  mass  NaN    NaN    NaN    NaN    NaN    NaN    NaN
```

```
f7    f8    f9 ...    f18    f19    f20  f0  0.000026  0.000924  0.039924 ...
0.119041  0.000152  0.000617  f1  0.315357  0.000580  0.000141 ... 0.000079  0.195312
0.000254  f2  0.000024  0.232319  0.000426 ... 0.000026  0.000399  0.092043  f3
0.000091  0.000188  0.021872 ... 0.174068  0.000328  0.000047  f4  0.000435  0.233003
0.000084 ... 0.000338  0.000400  0.068271  f5  0.000740  0.000671  0.143509 ...
0.517752  0.000337  0.000305  f6  0.000491  0.000699  0.038961 ... 0.390014  0.000226
0.000303  f7    NaN  0.000887  0.000328 ... 0.000251  0.299012  0.000276  f8    NaN
NaN  0.000548 ... 0.000235  0.000124  0.051674  f9    NaN    NaN    NaN ... 0.146027
0.000023  0.000175  f10    NaN    NaN    NaN ... 0.497812  0.000294  0.000334  f11
NaN    NaN    NaN ... 0.000089  0.239606  0.000590  f12    NaN    NaN    NaN ...
0.000126  0.000859  0.054172  f13    NaN    NaN    NaN ... 0.148374  0.000217
0.000031  f14    NaN    NaN    NaN ... 0.714613  0.000295  0.000070  f15    NaN
NaN    NaN ... 0.000265  0.221685  0.000307  f16    NaN    NaN    NaN ... 0.000242
0.000065  0.055865  f17    NaN    NaN    NaN ... 0.005241  0.000265  0.000239  f18
NaN    NaN    NaN ...    NaN  0.000117  0.000113  f19    NaN    NaN    NaN ...
NaN    NaN  0.000138  f20    NaN    NaN    NaN ...    NaN    NaN    NaN  f21
NaN    NaN    NaN ...    NaN    NaN    NaN  f22    NaN    NaN    NaN ...    NaN
NaN    NaN  f23    NaN    NaN    NaN ...    NaN    NaN    NaN  f24    NaN    NaN
NaN ...    NaN    NaN    NaN  f25    NaN    NaN    NaN ...    NaN    NaN    NaN
f26    NaN    NaN    NaN ...    NaN    NaN    NaN  mass    NaN    NaN    NaN ...
NaN    NaN    NaN
```

```

f21  f22  f23  f24  f25  f26  m f0  0.034794 0.080169 0.109947 0.289670
0.296976 0.327533 0.126 f1  0.000191 0.000100 0.000384 0.000178 0.000563 0.000782
0.000 f2  0.000312 0.000171 0.000075 0.000499 0.000186 0.000315 0.000 f3  0.010787
0.096627 0.190014 0.056976 0.406295 0.482503 0.203 f4  0.000485 0.000136 0.000061
0.000009 0.000346 0.000414 0.000 f5  0.044609 0.195117 0.188721 0.050882 0.082941
0.308311 0.064 f6  0.034123 0.209496 0.433089 0.126159 0.581645 0.835995 0.310 f7
0.000131 0.000023 0.000133 0.000182 0.000020 0.000621 0.000 f8  0.000047 0.000099
0.000354 0.000657 0.000375 0.000296 0.000 f9  0.307213 0.155551 0.015394 0.000711
0.235398 0.049122 0.055 f10 0.024445 0.211504 0.441256 0.123966 0.407570 0.760416
0.245 f11 0.000412 0.000509 0.000661 0.000559 0.001009 0.000406 0.000 f12 0.000263
0.000215 0.000447 0.000077 0.000146 0.000387 0.000 f13 0.306429 0.103286 0.000119
0.018679 0.055174 0.066454 0.000 f14 0.039548 0.234758 0.380652 0.104323 0.279477
0.656016 0.174 f15 0.000801 0.000806 0.001035 0.000278 0.000507 0.000914 0.000 f16
0.000275 0.000115 0.000236 0.000256 0.000441 0.000048 0.000 f17 0.205639 0.076665
0.033371 0.001520 0.076047 0.009786 0.028 f18 0.144685 0.184386 0.235264 0.082366
0.176284 0.521591 0.117 f19 0.000133 0.000185 0.000421 0.000060 0.000246 0.000005
0.000 f20 0.000442 0.000109 0.000493 0.000050 0.000220 0.000070 0.000 f21  NaN
0.153676 0.037039 0.003852 0.077928 0.021282 0.028 f22  NaN  NaN 0.727996
0.039866 0.155997 0.261498 0.057 f23  NaN  NaN  NaN 0.064041 0.302942
0.461259 0.155 f24  NaN  NaN  NaN  NaN 0.169888 0.141599 0.051 f25  NaN
NaN  NaN  NaN  NaN 0.572478 0.253 f26  NaN  NaN  NaN  NaN  NaN
NaN 0.323 mass  NaN  NaN  NaN  NaN  NaN  NaN

```

[28 rows x 28 columns]

```
In [16]: to_drop = [column for column in upper_tri.columns if any(upper_tri[column]
print((to_drop))
```

[]

```
In [17]: df = pd.get_dummies(df, columns=['f9', 'f13', 'f17', 'f21'], prefix=['f
```

```
In [18]: x = df.drop(['# label'],axis=1) ind_columns =
df.drop(['# label'],axis=1).columns y = df['# label']
```

We did normalize the attributes using StandardScaler() to scale them between 0 and 1 before running models.

In [19]:

```
# Normalize the data
scaler = StandardScaler ()
X_scaled = scaler .fit_transform (X)
```

In [20]:

```
from sklearn.model_selection import train_test_split
#Direct train/test split
X_train , X_test , Y_train , Y_test = train_test_split (
    X_scaled , y, test_size =0.20 , random_state =1234
)
```

Neural Network Model

In [21]:

```
import tensorflow as tf
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.optimizers import Adam
```

In [22]:

```
my_model = tf . keras . Model()
```

Metal device set to: Apple M1 Pro

2022-03-25 16:26:02.223468: I tensorflow/core/common_runtime/pluggable_

2022-03-25 16:26:02.223589: I tensorflow/core/common_runtime/pluggable_

Model 1 gelu - Four layers

In [23]:

```
layer_zero = tf.keras.Input(shape=(32,)) layer1 = tf.keras.layers.Dense(1094,
activation='gelu')(layer_zero) layer2 = tf.keras.layers.Dense(512,
activation='gelu')(layer1) layer3 = tf.keras.layers.Dense(128,
activation='gelu')(layer2) layer4 = tf.keras.layers.Dense(2,
activation='sigmoid')(layer2)
my_model = tf.keras.Model(inputs=layer_zero, outputs=layer4)
```

In [24]:

```
my_model.compile(optimizer=Adam(lr=1e-2), loss=tf.keras.losses.SparseCa safety
= EarlyStopping(monitor='val_loss', patience=3, min_delta=2e-4)
```

In [25]:

```
x_train, x_val, y_train, y_val = train_test_split(X_train, Y_train, tes
nn_history = my_model.fit(x_train, y_train,
validation_data=(x_val, y_val),
callbacks=[safety], epochs=50,
batch_size=100)
```

Epoch 1/50

2022-03-25 16:26:04.870768: W tensorflow/core/platform/profile_utils/cp

```
2022-03-25 16:26:05.102127: I tensorflow/core/grappler/optimizers/custo
```

```
2188/2188 [=====] - ETA: 0s - loss: 0.2856 - a
```

```
2022-03-25 16:26:34.969380: I tensorflow/core/grappler/optimizers/custo
```

```
2188/2188 [=====] - 33s 15ms/step - loss: 0.28 Epoch 2/50
```

```
2188/2188 [=====] - 33s 15ms/step - loss: 0.27 Epoch 3/50
```

```
2188/2188 [=====] - 31s 14ms/step - loss: 0.26 Epoch 4/50
```

```
2188/2188 [=====] - 30s 14ms/step - loss: 0.26 Epoch 5/50
```

```
2188/2188 [=====] - 30s 14ms/step - loss: 0.26 Epoch 6/50
```

```
2188/2188 [=====] - 31s 14ms/step - loss: 0.26 Epoch 7/50
```

```
2188/2188 [=====] - 31s 14ms/step - loss: 0.26 Epoch 8/50
```

In [26]:

```
2188/2188 [=====] - 31s 14ms/step - loss : 0.26
```

```
my_model.summary()
```

```
Model: "model_1"
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 32)]	0
dense (Dense)	(None, 1094)	36102
dense_1 (Dense)	(None, 512)	560640
dense_3 (Dense)	(None, 2)	1026
		===== Total
params: 597,768		
Trainable params: 597,768		
Non-trainable params: 0		



In [27]:

```
history_df = pd.DataFrame(nn_history.history) history_df[['loss',
'val_loss']].plot()

history_df = pd.DataFrame(nn_history.history) history_df[['accuracy',
'val_accuracy']].plot()
```

Out[27]:

```
<AxesSubplot:>
```



Model 1 gelu - 3 layers

```
In [53]: layer_zero = tf.keras.Input(shape=(32,)) layer1 = tf.keras.layers.Dense(512,
activation='gelu')(layer_zero) layer2 = tf.keras.layers.Dense(256,
activation='gelu')(layer1) #layer3 =
tf.keras.layers.Dense(128,
activation='gelu')(layer2) layer4 = tf.keras.layers.Dense(2,
activation='sigmoid')(layer2)
my_model = tf.keras.Model(inputs=layer_zero, outputs=layer4)
```

```
In [54]: my_model.compile(optimizer=Adam(lr=1e-2), loss=tf.keras.losses.SparseCa safety
= EarlyStopping(monitor='val_loss', patience=3, min_delta=2e-4)
x_train, x_val, y_train, y_val = train_test_split(X_train, Y_train, tes
```

```
In [55]: nn_history = my_model.fit(x_train, y_train,
validation_data=(x_val, y_val),
callbacks=[safety], epochs=50,
batch_size=100)

Epoch 1/50
2022-03-26 10:27:51.674809: I tensorflow/core/grappler/optimizers/custo
44800/44800 [=====] - ETA: 0s - loss: 0.2929 -
2022-03-26 10:31:34.644074: I tensorflow/core/grappler/optimizers/custo
44800/44800 [=====] - 255s 6ms/step - loss: 0 Epoch 2/50
44800/44800 [=====] - 249s 6ms/step - loss: 0 Epoch 3/50
44800/44800 [=====] - 251s 6ms/step - loss: 0 Epoch 4/50
44800/44800 [=====] - 253s 6ms/step - loss: 0 Epoch 5/50
44800/44800 [=====] - 245s 5ms/step - loss: 0 Epoch 6/50
44800/44800 [=====] - 244s 5ms/step - loss: 0


In [56]: y_pred=my_model.predict(X_test)

cm = confusion_matrix(Y_test, np.argmax(y_pred, axis=1))

cm_plot_label=['No', 'Yes'] plot_confusion_matrix(cm, cm_plot_label, title
='Confusion Matrix')

2022-03-26 10:52:50.264679: I tensorflow/core/grappler/optimizers/custo

[[601100 98345]
 [ 74978 625577]]
```



In [57]:

```
my_model.summary()
```

Model: "model_5"

Layer (type)	Output Shape	Param #
input_5 (InputLayer)	[(None, 32)]	0
dense_13 (Dense)	(None, 512)	16896
dense_14 (Dense)	(None, 256)	131328
dense_15 (Dense)	(None, 2)	514

Total params: 148,738
 Trainable params: 148,738
 Non-trainable params: 0

In [32]:

```



history_df = pd.DataFrame(nn_history.history) history_df[['loss',
'val_loss']].plot()

history_df = pd.DataFrame(nn_history.history) history_df[['accuracy',
'val_accuracy']].plot()

```

Out[32]:

<AxesSubplot:>



Batch size 1024

In [33]:

```
nn_history = my_model.fit(x_train, y_train,
validation_data=(x_val, y_val),
callbacks=[safety],          epochs=50,
batch_size=1024)
```

Epoch 1/50

2022-03-25 16:35:50.010830: I tensorflow/core/grappler/optimizers/custo

```
4375/4375 [=====] - 29s 7ms/step - loss: 0.267 Epoch 2/50
4375/4375 [=====] - 31s 7ms/step - loss: 0.266 Epoch 3/50
4375/4375 [=====] - 29s 7ms/step - loss: 0.266 Epoch 4/50
4375/4375 [=====] - 29s 7ms/step - loss: 0.266 Epoch 5/50
4375/4375 [=====] - 29s 7ms/step - loss: 0.266 Epoch 6/50
4375/4375 [=====] - 27s 6ms/step - loss: 0.266 Epoch 7/50
```

In [34]:

4375/4375 [=====] - 28s 6ms/step -

loss : 0.265

my_model.summary()

Model: "model_2"

Layer (type)	Output Shape	Param #
=====		
input_2 (InputLayer)	[(None, 32)]	0
dense_4 (Dense)	(None, 512)	16896
dense_5 (Dense)	(None, 256)	131328
dense_6 (Dense)	(None, 2)	514
=====		
Total params: 148,738		
Trainable params: 148,738		
Non-trainable params: 0		

In []:


In [35]:


```
history_df = pd.DataFrame(nn_history.history) history_df[['loss',
'val_loss']].plot()
```

```
history_df = pd.DataFrame(nn_history.history) history_df[['accuracy',
'val_accuracy']].plot()
```

Out[35]:

<AxesSubplot:>





Model 2 Swish - 3 layers

```
In [36]: layer_zero = tf.keras.Input(shape=(32,)) layer1 = tf.keras.layers.Dense(512,
activation='swish')(layer_zero) layer2 = tf.keras.layers.Dense(256,
activation='swish')(layer1) #layer3 =
tf.keras.layers.Dense(128,
activation='gelu')(layer2) layer4 = tf.keras.layers.Dense(2,
activation='sigmoid')(layer2)
my_model = tf.keras.Model(inputs=layer_zero, outputs=layer4)
```

```
In [37]: my_model.compile(optimizer=Adam(lr=1e-2), loss=tf.keras.losses.SparseCa safety
= EarlyStopping(monitor='val_loss', patience=3, min_delta=2e-4)
x_train, x_val, y_train, y_val = train_test_split(X_train, Y_train, tes
```

```
In [38]: nn_history = my_model.fit(x_train, y_train,
validation_data=(x_val, y_val),
callbacks=[safety], epochs=50,
batch_size=2048)
```

Epoch 1/50

2022-03-25 16:39:18.261506: I tensorflow/core/grappler/optimizers/custo

2188/2188 [=====] - ETA: 0s - loss: 0.2841 - a

2022-03-25 16:39:33.405411: I tensorflow/core/grappler/optimizers/custo

2188/2188 [=====] - 17s 8ms/step - loss: 0.284 Epoch 2/50

2188/2188 [=====] - 16s 7ms/step - loss: 0.271 Epoch 3/50

2188/2188 [=====] - 16s 7ms/step - loss: 0.269 Epoch 4/50

2188/2188 [=====] - 16s 7ms/step - loss: 0.268 Epoch 5/50

2188/2188 [=====] - 16s 8ms/step - loss: 0.267 Epoch 6/50

2188/2188 [=====] - 16s 7ms/step - loss: 0.266 Epoch 7/50

```

2188/2188 [=====] - 17s 8ms/step - loss: 0.265 Epoch 8/50
2188/2188 [=====] - 16s 7ms/step - loss: 0.264 Epoch 9/50
2188/2188 [=====] - 16s 7ms/step - loss: 0.265 Epoch 10/50
2188/2188 [=====] - 16s 7ms/step - loss: 0.264 Epoch 11/50
2188/2188 [=====] - 16s 7ms/step - loss: 0.264

```

In [39]:

```
my_model.summary()
```

Model: "model_3"


Layer (type)	Output Shape	Param #
input_3 (InputLayer)	[(None, 32)]	0
dense_7 (Dense)	(None, 512)	16896
dense_8 (Dense)	(None, 256)	131328
dense_9 (Dense)	(None, 2)	514
Total params: 148,738		
Trainable params: 148,738		
Non-trainable params: 0		

In [40]:

```
history_df = pd.DataFrame(nn_history.history) history_df[['loss',
'val_loss']].plot()

history_df = pd.DataFrame(nn_history.history) history_df[['accuracy',
'val_accuracy']].plot()
```

Out[40]:



In [41]:

```
my_model.summary()
```

Model: "model_3"

Layer (type)	Output Shape	Param #
=====		
input_3 (InputLayer)	[(None, 32)]	0
dense_7 (Dense)	(None, 512)	16896
dense_8 (Dense)	(None, 256)	131328
dense_9 (Dense)	(None, 2)	514
=====		

Total params: 148,738
 Trainable params: 148,738
 Non-trainable params: 0

Model Evaluation

In [42]: `y_pred =my_model.predict (X_test)`
 2022-03-25 16:42:18.676823: I tensorflow/core/grappler/optimizers/custo

In [43]: `y_pred`

Out[43]: array([[0.48715827, 0.46217152],
 [0.08716998, 0.96522045],
 [0.5221664 , 0.43013135], ...,
 [0.23122704, 0.9607484],
 [0.17528357, 0.8504923],
 [0.24064562, 0.7891484]], dtype=float32)

In [44]: `from sklearn.metrics import classification_report`
`print(classification_report(Y_test, np.argmax(y_pred, axis=1)))`

```

          precision  recall f1-score  support
0.0    0.90    0.85    0.88   699445
1.0    0.86    0.91    0.88   700555
accuracy              0.88  1400000  macro avg    0.88
0.88   0.88  1400000 weighted avg    0.88   0.88   0.88
1400000

```

In [47]: `from sklearn.metrics import`
`confusion_matrix import itertools`

`def plot_confusion_matrix(cm, classes,`
 `title='Confusion matrix', cmap=plt.cm.Blues):`

`print(cm)`

`plt.imshow(cm, interpolation='nearest', cmap=cmap)`
 `plt.title(title) plt.colorbar()`
 `tick_marks = np.arange(len(classes)) plt.xticks(tick_marks, classes,`
`rotation=55) plt.yticks(tick_marks, classes)`
 `#fmt =`
 `'.2f' 'd' thresh =`
`cm.max() / 2.`

```

for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, format(cm[i, j]), horizontalalignment="center",
             color="white" if cm[i, j] > thresh else "black")

plt.ylabel('True label') plt.xlabel('Predicted label')
plt.tight_layout()

cm = confusion_matrix(Y_test, np.argmax(y_pred, axis=1))


cm_plot_label = ['No', 'Yes']
plot_confusion_matrix(cm, cm_plot_label, title='Confusion Matrix')

```

```

[[597416 102029]
 [ 64132 636423]]

```




In [46]:

```

from sklearn.metrics import roc_auc_score , auc
from sklearn.metrics import roc_curve
roc_log = roc_auc_score ( Y_test , np.argmax(y_pred , axis =1))
false_positive_rate , true_positive_rate , threshold = roc_curve ( Y_test ,
area_under_curve = auc( false_positive_rate , true_positive_rate )

plt . plot ([ 0, 1], [ 0, 1], 'r--' )
plt . plot ( false_positive_rate , true_positive_rate , label ='AUC = {:.3f}'.
plt . xlabel ( 'False positive rate' )
plt . ylabel ( 'True positive rate' )
plt . title ( 'ROC curve' )
plt . legend ( loc ='best' )
plt . show()
plt . close ()

```



In []:

In []:

In []: